

## Chapter 4 - Balancing the Gathering and Use of Information

→ This chapter is about learning from evaluative feedback.

What is evaluative feedback? (Recap)

→ It is a summary for the learner of how well he or she has performed on a particular task. In terms of RL, the reward signal provides evaluative feedback to the agent.

→ In general, every decision made by an intelligent entity is a trade-off between exploration and exploitation.

→ It is important for the agent to "correctly" decide when to gather more knowledge from the environment and when to take an action based on previously learned knowledge.

Exploration builds the knowledge that allows for effective exploitation, and maximum exploitation is the goal of any decision maker.

Def<sup>n</sup>:

Horizon: The number of time steps that the agent interacts with the environment in each episode.

Greedy Horizon: A special task with horizon 1, i.e. each episode is exactly 1 time step.

Multi-Armed Bandits: (MAB's)

- Special kind of RL problem in which the environment's MDP has a single state, multiple actions and each episode only lasts 1 time step (in other words MAB's have a greedy horizon).
- Ex: Youtube trying to show you an ad, e-commerce websites recommending products.
- MAB's highlight the exploration - exploitation tradeoff, so their study can help agents in more complex environments also.

Setup

- The goal of our task is to maximise the expected total reward in some 'n' time steps, (in other words learn the p.d of reward signal as accurately as possible)
- Each action has an expected reward. Call this the value of the action. (Keep in mind that the agent does not know the value of each action)
- From past experience, the agent may have decent estimates of the value of each action. In a given episode, if the agent selects the action with highest reward (by its estimate), then this action is "greedy".
- Greedy actions correspond to exploiting and on the other hand if the agent chooses an action with lower reward estimate, then in that time step, it is exploring.
- Exploitation is the right thing to do to maximise reward in a single time step, but exploration might yield higher returns in the long run.



Recall the action-value function :

$$q_{\pi}(s, a) = \mathbb{E}_{\pi} [G_t | S_t = s, A_t = a]$$

In the MAB setting, as there is only a single state,

$$q(a) = \mathbb{E}[R_t | A_t = a]$$

→ Here,  $V_*$  (optimal state-value function) will represent the highest expected reward we can get in a single time step :

$$V_* = q(a_*), \text{ where } a_* = \arg \max_{a \in A} q(a)$$

Regret :

→ we want to tweak the statement of our goal which is "maximize the expected reward".

→ Why? Agent A may find the optimal action in the final episode after many episodes of random actions, while another Agent B might be much more efficient and find the optimal action much quicker than A. If our goal was "maximize the expected reward", both A and B would be thought to have achieved the goal, but obviously we want agent B.

So, a more robust way to capture a more complete goal is for the agent to maximize the per-episode expected reward, while still minimizing the 'total expected reward loss of rewards across all episodes'.

→ We will call this quantity **total regret**, and to calculate it, we sum the per-episode difference of the true expected reward of the optimal action and the true expected reward of the selected action.

→ True expected reward, requires access to the environment's MDP (the action value function  $q$  in MAB's), to be calculated.

(Total Regret will only be calculated to compare the behaviour of 2 agents)

$$\text{Total Regret } T = \sum_{e=1}^E \mathbb{E}[v_* - q(A_e)]$$

## Approaches to Solving MAB's

3 major kinds:

- 1) Random Exploration Strategies - Agent will mostly exploit but randomly explore.
- 2) Optimistic Exploration Strategies Quantifies uncertainty, and increases preference for states with highest uncertainty.
- 3) Information state - space exploration Strategies: Model the information state of the agent into the environment, i.e. tries to encode the uncertainty of the agent into the ~~environment~~ state space. Risks large state spaces and high complexity.

→ Note that in any strategy, the estimation of the action-value function  $Q$ , is done in the same way:

If action  $a$  is performed in  $k$  time steps/episodes and yields reward  $Q_i(a)$ , then our estimate of  $Q(a)$  will be

$$Q(a) \approx \sum_{i=1}^k Q_i(a) / k$$

→ So strategies differ in how the estimate is used, not how the estimation is performed.



## Two example Strategies:

### 1) Greedy (Always Exploit) :

- Always select action with highest estimated  $Q$  value.
- There is some chance, that the action we select first is optimal, but this chance decreases considerably as action space grows larger.
- In an environment with  $Q$ -values <sup>(estimates)</sup> initialized to zero, and no negative rewards, our agent will get stuck with the first action it chooses.
- An agent that is always exploiting, will often fail to find the optimal action, however, in general, if an agent only has a limited number of 'training' episodes left, greedy strategy is effective.
- The agent is 'short-sighted' if it is always exploiting as it does not trade-in immediate satisfaction for information gain that could lead to higher future rewards.

## 2) Random (Always Explore):

- Never exploits and sole goal of agent is to gain information.
- This strategy is also very bad, and in case the agent does 'explore' the optimal action, it ~~will~~<sup>might</sup> never come back to it as it is always exploring.
- We want algorithms to balance both exploration and exploitation.
- Note that random exploration is just one way to explore various actions. Other pure exploration strategies can also be designed.

## Some viable Strategies:

### 1) $\epsilon$ -Greedy (Almost always greedy and sometimes random.)

- Select action agent thinks is best almost all the time, but randomly explore other actions as well.
- Agent's estimate of action-value function has an opportunity to converge to  $Q$  (true action-value function)



Sample  $\epsilon$ -greedy Algo: ~~Q~~ ( $Q$  is lookup table with  $q$ -value estimates)

$\epsilon$ -greedy (MAB,  $\epsilon = e$ , episodes =  $n$ ) : ( $0 < e < 1$ )

1 Sample a random number in  $(0, 1)$  (we sample from uniform distribution on  $[0, 1]$ )

if  $\text{rand} > e$ :

# exploit

$\text{action} = \text{argmax}(Q)$

else:

# explore

sample a random integer in  $\text{len } Q$  ( $\text{randint}$ )

$\text{action} = \text{randint}$

Remarks:  $\rightarrow$  This algo has randomness at both stages:

1) Choosing between exploitation and exploration

2) Exploration.

$\rightarrow$  Probability of exploration is exactly  $e$  and exploitation is  $1-e$ .

$\rightarrow$  Note that in the exploration stage, greedy action may also be picked, so in reality, exploration probability ~~might be~~ is even less than ' $e$ '.

- 2) Decaying  $\epsilon$  Greedy: (First maximize exploration, then exploit)
- Explore more initially when agent hasn't experienced the environment much.
  - As value function estimates improve, exploit more.
  - Simple way to do this: Start with some  $0 < \epsilon < 1$  and decay its value at every step.

Sample algo: Linear decaying  $\epsilon$ -greedy.

lindec-eps-greedy (MAB,  $\epsilon_i$ ,  $\epsilon_m$ ,  $\gamma$ , episodes = n) :

( $\epsilon_i$  is initial  $\epsilon$ ,  $\epsilon_m$  is minimum  $\epsilon$ ,  $\gamma$  is decay ratio)

for  $i$  in range(n) :

    decay\_episodes = episodes \*  $\gamma$

    # Calculate  $\epsilon$  value for current episode

$\epsilon = 1 - i / \text{decay\_episodes}$

$\epsilon += \epsilon_m$

$\epsilon += \epsilon_m$

↘ Now generate random number and proceed as in  $\epsilon$  greedy.

Exponential  $\epsilon$  decay can also be executed instead of linear.

→ In general, many ways to execute  $\epsilon$  decay, but all of them should favour exploration in the beginning.

### 3) Optimistic Initialization

→ Treats actions that were not sufficiently explored as the best actions.

→ Initialize agent's  $Q$  function estimate to a high value and then start acting greedily, using these estimates.

(Usually good high value is not available, but assume that for now we do)

→ Counts of actions performed can act as a confidence measure in the agent's estimate of the action's  $Q$  value. So, initializing counts to high values is also necessary, otherwise  $Q$  function will change too quickly.

→ Agent initially expects high rewards, but as experience is gained,  $Q$  value estimates drop and start to converge.

→ Algo is straight forward : Initialization + pure exploitation.



## More advanced Strategies:

1) Softmax (selecting actions randomly in proportion to their estimates)

→ It is a random exploration strategy that takes into account estimated Q values to make a decision.

→ If an action has a low estimate, agent will try it with less likelihood.

Softmax Strategy does ~~exactly~~ exactly this:

→ Probability of selecting an action  $\propto$  current Q value estimate of that action.

→ Hyperparameter,  $T$  (called temperature) can also be added to control the algo's sensitivity to Q value estimates.

→  $T$  can also be decayed.

→ Let  $\pi$  denote the policy, i.e. returns probability of an action being taken.  $\pi$  depends on current Q ~~est~~ function

estimate: 
$$\pi(a) = \frac{e^{(Q(a)/T)}}{\sum_{i=1}^K e^{(Q(a_i)/T)}}$$

→ As  $T \uparrow$ ,  $\pi$  resembles uniform distribution.

## UCB:

- First, some issues with optimistic initialization:
  - we don't know maximum reward agent can obtain, so setting initial values ~~low~~ much higher than this would lead to huge training times.
  - If initial values are much smaller than actual max reward, then agent will probably never discover optimal action.
  - Counts were also initialized, but we want to be able to use counts to quantify uncertainty.
- The Upper Confidence Bound Strategy (UCB) is an improvement, in some sense, to optimistic initialization.
- In UCB, uncertainty of value estimates is taken into consideration.

$$A_e = \underset{a}{\operatorname{argmax}} \left[ \underset{\substack{\downarrow \\ \text{Q value estimate}}}{Q_e(a)} + c \sqrt{\frac{\ln e}{N_e(a)}} \right] \rightarrow \text{Uncertainty Bonus}$$

$\downarrow$   
action to be taken at episode  $e$

$\downarrow$   
# of times  $a$  was performed until episode ' $e$ '.

### 3) Thompson Sampling


→ Keep track of each Q-value estimate as a Gaussian distribution (or any other suitable distribution).

→ In case of Gaussian distribution, there are 2 parameters  $\mu$  and  $\sigma$ .  $\mu$  will be the current Q value estimate, while s.d  $\sigma$ , will measure the uncertainty of the ~~episode~~ estimate.

→  $\mu$  and  $\sigma$  are updated on each episode.

→ Whenever action needs to be taken, sample from these distributions and pick the action with highest sample.

→ As number of episodes  $\uparrow$ , means are supposed to get closer and closer to actual q value and  $\sigma$ 's should drop.



→ Progression of the distribution of some action as episodes increase.